



From Data Sharing to Clean Rooms

Exploring Collaboration Options in Snowflake

Tom Manfredi
Senior Partner Sales Engineer
January 2025

What you will learn...

- Common governance challenges
- Fundamentals of Secure Data Sharing in Snowflake
- Spectrum of Privacy Options in Snowflake Collaboration
 - Simple Sharing w/RBAC
 - Privacy with Views
 - Dynamic/Conditional Data Masking
 - Tag-based Policies and Propagation
 - Row Access Policies
 - Projection Policies
 - Differential Privacy
 - Data Clean Rooms
- How to best balance collaboration goals
 - Analytic Value / Privacy / Ease of Use



DATA GOVERNANCE CHALLENGES

Data Is Everywhere



Must be able to eliminate silos inside and outside your organization

Managing Data Is Unnecessarily Complex



Knowing what your data is — and how it is being used — is hard

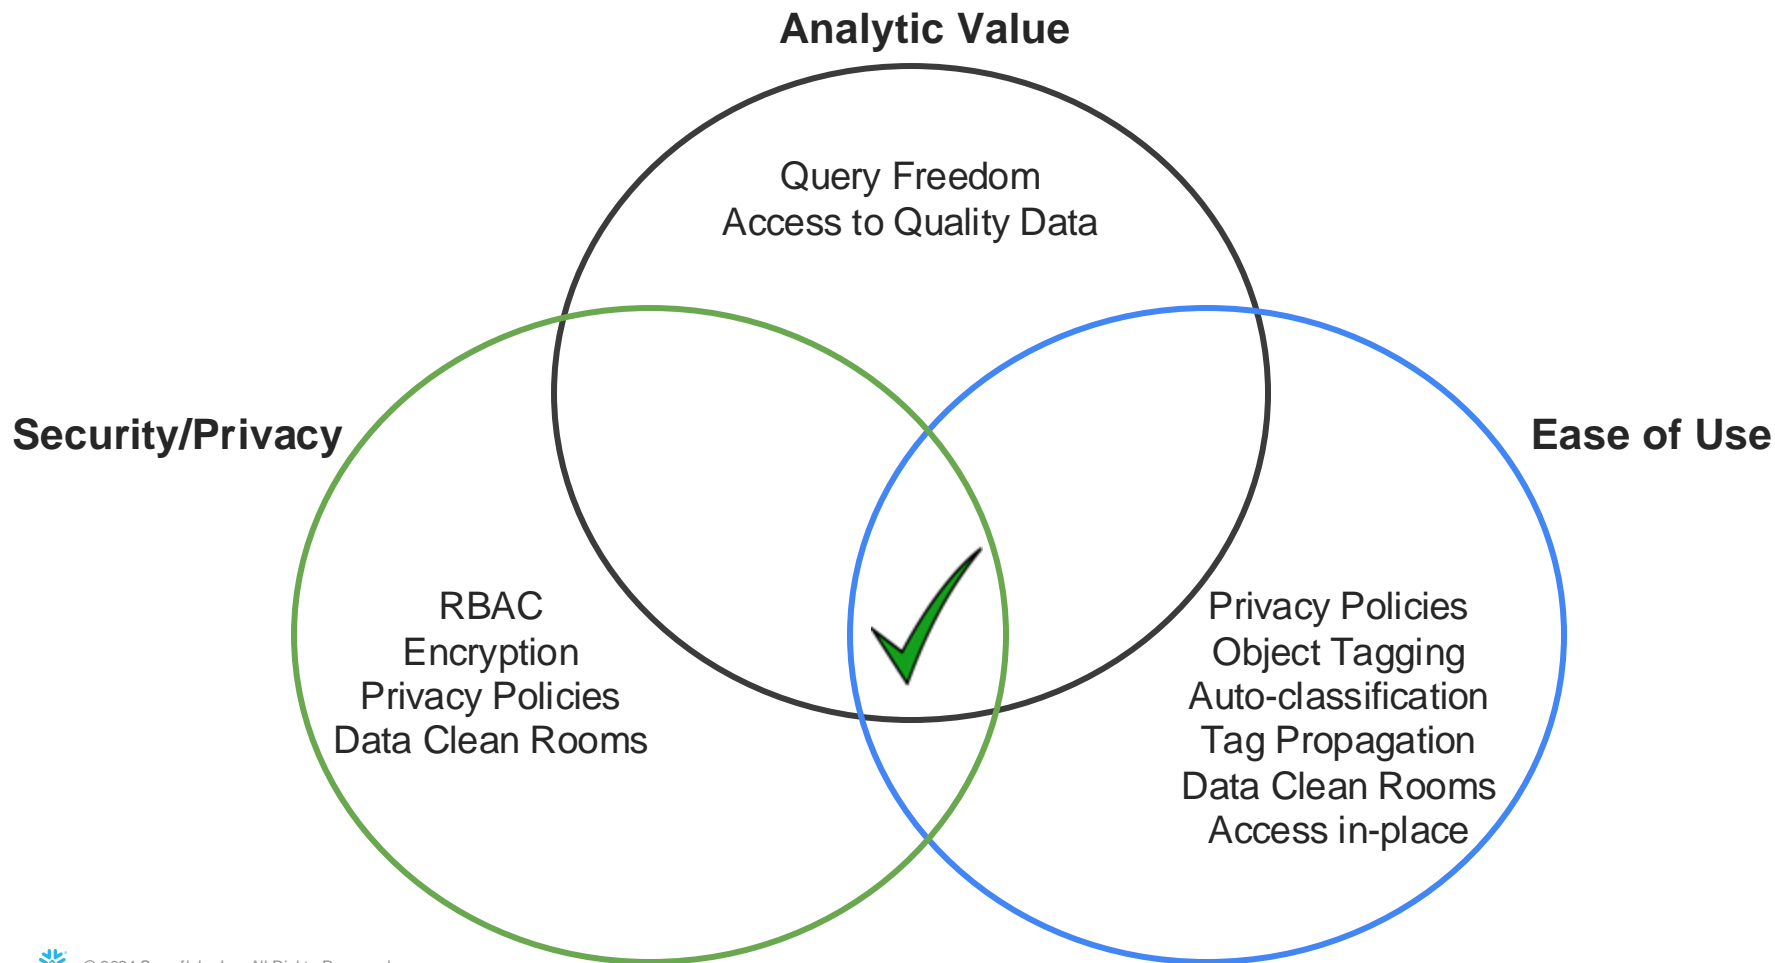
Security and Governance Are Inherently Rigid



Requires a flexible approach to managing risk, regardless of workload



Goal is to balance...



Snowflake Makes Collaboration Easy

Traditional Methods


FTP | APIs | ETL | Federated Protocols



- ✗ **Costly to maintain** data pipeline infrastructure to share beyond a single cloud region, delaying access and exposing governance risk
- ✗ **Heavy data wrangling** required to train and distribute AI Models
- ✗ Integrating SaaS requires copying data to the app; **creating new silos**

Snowflake

Privacy-Preserving Collaboration

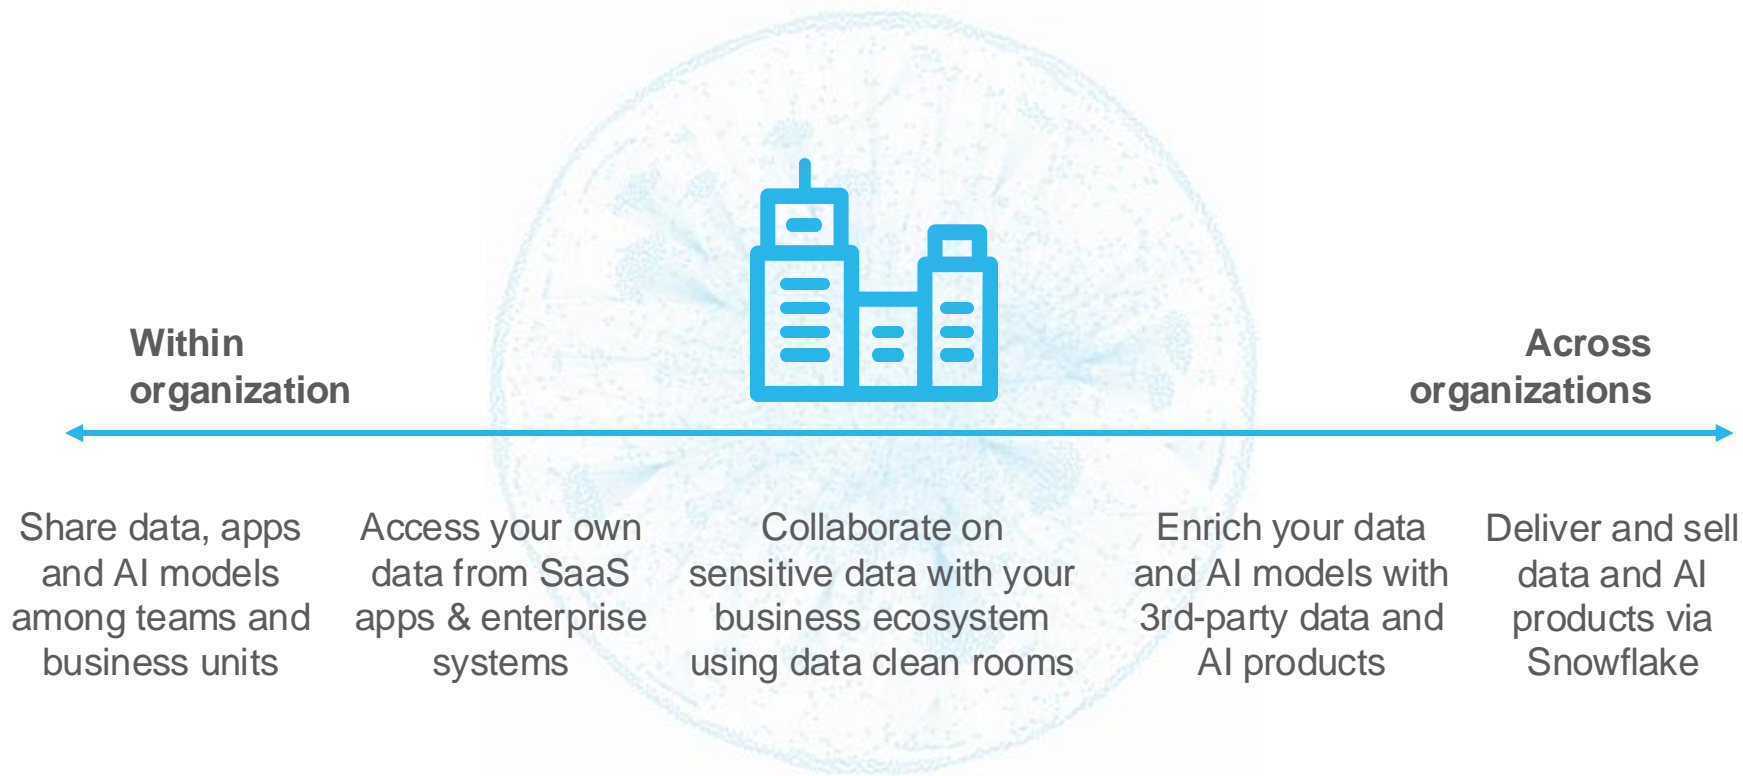


- ✓ Share a **single managed copy** of data across cloud region; no ETL or copies giving you fast delivery and control with insights derived from shared data
- ✓ Direct access to governed data for AI model training and **easy distribution** to other teams, partners and customers
- ✓ Securely install apps directly into your Snowflake account to **bring the code to your data**



Collaboration in The AI Data Cloud

Discover, Share & Monetize Data, Apps and AI Products Across Clouds



We start with a table...

Customer

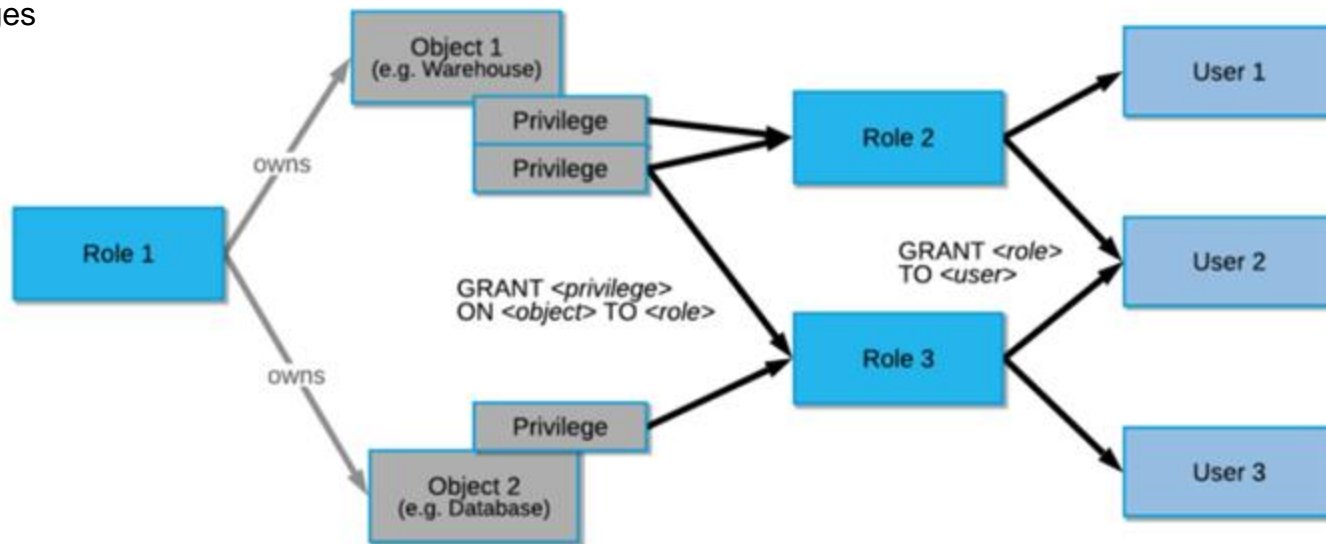
name	gender	age	zip_code	phone
John Smith	male	39	79007	123-555-1234
Jane Doe	female	50	77001	333-555-1236
Mary Taylor	female	46	77020	222-333-1111
Gene Marshall	non-binary	48	77042	555-555-1234
Michael Gaines	male	75	79003	666-666-1357

- The table has been instantiated from the encrypted, at-rest files (micro-partitions)
- The information in the table is opaque to Snowflake
- “Clear text” data is only visible to
 - Authenticated users to Snowflake
 - Assigned to an authorized role



Role-based Access Control (RBAC)

- Objects (data, applications, models...)
- Privileges
- Roles
- Users



Access privileges are assigned to roles, which are in turn assigned to users.

RBAC can help us down to the table/view level. Beyond that, we need something else



DEMO




Data Clean Rooms



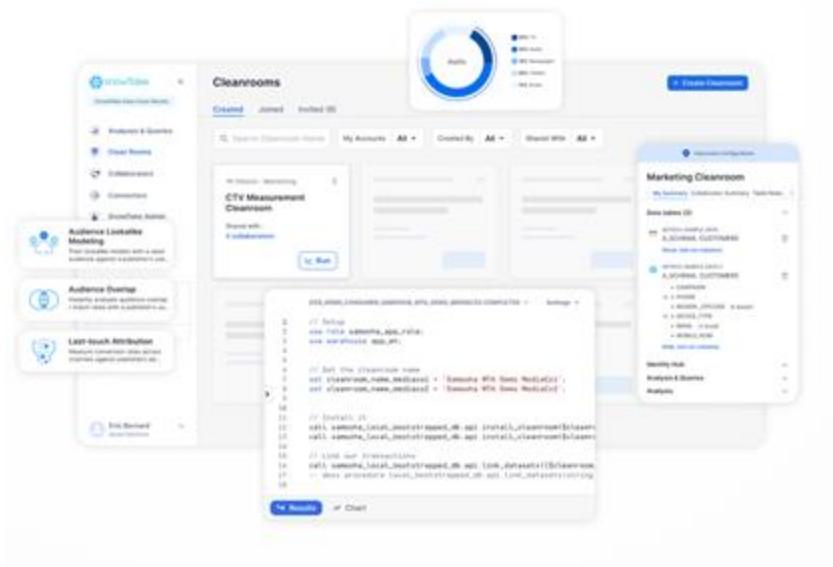
Data Clean Rooms

A secure environment where multiple parties can collaborate on sensitive or regulated data without exposing or moving the underlying data



Snowflake Data Clean Rooms

An app that makes data clean rooms easy to use & customizable for non-technical and technical users respectively.



- Allows for **secure collaboration on sensitive data** without exposing the underlying data
- Requires *no data movement*
- Snowflake Native App **available for free** on Snowflake Marketplace
- Enables **business users** with **ready-to-use industry-specific templates**
- **Customizable templates** using developer APIs



Snowflake Differentiators



Native
Applications



Secure
Data Sharing



Snowpark



Row Access
Policies



Stored
Procedures



Snowgrid



Conclusions: Balancing

Analytic Value
Security/Privacy
Ease of Use

Secure Data Sharing

+

Secure Views

+

Policy Constrained access

Row Access Policies
Dynamic Data Masking
Conditional Masking
Projection Policies
Aggregation Policies
Differential Privacy

+

Pre-Defined Access (templated)

Custom Native App
Data Clean Room

Easy to setup and maintain
Provider can **audit** and **monitor** usage



Provides **conditional** control over attributes (**columns**) and **rows** exposed to audience



Flexible control over how data can be viewed (**projected**) and **aggregated** without pre-defined access (**templates**)

Advanced setup and usage based on pre-defined templates



The audience is trusted to be able to access and use data “freely”

All audiences

No protection against (non friendly) sophisticated inspection

Internal & trusted audience

Protects against misuse or accidental misuse of the data - keeps usage within established boundaries

“Friendly” audience

Protects against complex Data Clean Room “attacks”

External, semi- or untrusted audience

Ease of use

Easy & flexible

more complex

Trusted

untrusted

Level of trust Provider/Consumer





THANK YOU



Backup Slides




RBAC, DAC, Views, UDFs

RBAC & DAC

Views & UDFs

name	age	zip_code
J Smith	3-	790**
J Doe	5-	770**
M Taylor	4-	770**
M Gaines	7-	790**

name	gender	age	zip_code	phone
John Smith	male	39	79007	123-555-1234
Jane Doe	female	50	77001	333-555-1236
Mary Taylor	trans-fem	46	77020	222-333-1111
Gene Marshall	non-binary	48	77042	555-555-1234
Michael Gaines	male	75	79003	666-666-1357



RBAC & DAC Protect the table

- Every object in Snowflake is subject to these controls, and they are at the whole-object level
- RBAC inheritance and other RBAC features apply
- The customer controls RBAC completely
- DAC (Discretionary Access Control) applies to the role that owns the object, unless the object is subject to Managed Schema Access

You may also create Views & UDFs

- These are mostly used to redact or transform rows, columns, or even cells, and create a new object
- The new object has RBAC and DAC controls



Column-Level Security / Row Access Policy

RBAC & DAC

Views & UDFs

Column-Level Security

Row Access Policy

Conditional Policy

We can use Policy controls for Columns and Rows

- Prevent View/UDF explosion
- Table/View owners and privileged users (such as ACCOUNTADMIN) unauthorized to data by default
- Ensure controls are applied in any context where the object's data is used

We get more ease of management

- Centrally manage policies
- Apply a single policy to multiple tables
- Built-in separation of duty: policy admins assign and users are subject to policy controls
- All application and use is fully audited

FAMILYNAME	GENDER	AGE	ZIPCODE	PHONE	OPTIN
Romero	F	56	31675	4282478462	Y
Polk	N	52	16971	7565697078	U
Gaines	F	61	28877	4210779144	Y
Gordon	M	36	85115	6474995638	N
Hammer	F	63	50587	5947720813	Y



Dynamic Data Masking

RBAC & DAC

Views & UDFs

Column-Level Security

Dynamic Data Masking

We can leverage Column-level Security to dynamically mask data at query time

- No change to the stored data
- Mask or partially mask using constant value, hash, and custom functions
- Unmask for authorized users only

FAMILYNAME	GENDER	AGE	ZIPCODE	PHONE	OPTIN
Romero	F	56	31675	4282478462	Y
Polk	N	52	16971	7565697078	U
Gaines	F	61	28877	4210779144	Y
Gordon	M	36	85115	6474995638	N
Hammer	F	63	50587	5947720813	Y

Query results

```
phone      name
***-**-5534 ** masked **
***-**-3564 ** masked **
***-**-9787 ** masked **
```



Alex
(Unauthorized)

Query results

```
phone      name
408-123-5534 ** masked **
510-335-3564 ** masked **
214-553-9787 ** masked **
```



Morgan
(Partially Authorized)

```
create or replace masking policy F00
as (val string) returns string ->
case
  when is_granted_to_invoker_role('SEECLEAR')
  then val
  when current_role('ONLYPART')
  then regexp_replace(val, '[0-9]', '*', 7)
  when is_role_in_session('CRYPTO')
  then decrypt_raw(val, keyval, IV, ...)
  when is_role_in_session('BESPOKE')
  then user_defined_Func(val, baz, ...)
  else '** masked **'
end;
```

```
alter table start_with_a_table modify column PHONE
set masking policy F00;
```



External Tokenization

RBAC & DAC

Views & UDFs

Column-Level Security

External Tokenization

FAMILYNAME	GENDER	AGE	ZIPCODE	PHONE	OPTIN
Romero	F	56	31675	awiufyaf873fg	Y
Polk	N	52	16971	78ybhsbbcvzd	U
Gaines	F	61	28877	984iuwrfgjffss	Y
Gordon	M	36	85115	ciudsjhciasudg	N
Hammer	F	63	50587	8347ryfgvgshf	Y

Example using policy:

```
create or replace masking policy BAR as (val
string) returns string ->
case
  when is_granted_to_invoker_role('SEETOKENS')
  then val
  when current_role('GETREAL')
  then detok_ext_func(val, current_user(),...)
  else '** masked **'
end;
```

Example using SQL outside policy:


```
SELECT detok_ext_func(T1.phone) AS REAL_PHONE
, T1.GENDER
, T2.ZIP
FROM T1
JOIN T2
ON T2.PHONE = T1.PHONE
```

Ingest protected (PII/PHI) data as Externally Tokenized

- Using tokenization provider functionality upstream from Snowflake

De-tokenize for authorized users at query time

- The tokenization provider is called using a Snowflake External Function to de-tokenize data
- For unauthorized users, third-party service is not called
- Can be used in policy or outside



Row Access Policy

RBAC & DAC

Views & UDFs

Column-Level Security

Row Access Policy

FAMILYNAME	GENDER	AGE	ZIPCODE	PHONE	OPTIN
Romero	F	56	31675	4282478462	Y
Polk	N	52	16971	7565697078	U
Gaines	F	61	28877	4210779144	Y
Gordon	M	36	85115	6474995638	N
Hammer	F	63	50587	5947720813	Y

Query results

zip_code name
31675 ** masked **



Alex
(Only 31675)

Query results

zip_code name
28877 ** masked **



Morgan
(Only 28877)

Filter rows at query time based on user role and lookup table

- Policy contains condition(s) to allow or filter out rows
- Policy is applied to one or more table, view, or external table in an account
- Dynamically generated predicate filters out rows the user is not authorized to see at query time
- Can be combined with other controls

```
create or replace row access policy F00
as (this_zip varchar) returns boolean ->
'all_seeing_role' = current_role()
or
exists (
select 1 from zip_mapping_table
where info_reader = current_role()
and zip_code = this_zip);
```



Conditional Policy

RBAC & DAC

Views & UDFs

Column-Level Security

Row Access Policy

Conditional Policy

FAMILYNAME	GENDER	AGE	ZIPCODE	PHONE	OPTIN
Romero	F	56	31675	4282478462	Y
Polk	N	52	16971	7565697078	U
Gaines	F	61	28877	4210779144	Y
Gordon	M	36	85115	6474995638	N
Hammer	F	63	50587	5947720813	Y

Query results

```

phone      name
***MASKED*** Polk
4210779144 Gaines
***MASKED*** Gordon
  
```



All Users

Conditional policy is a type of Column-level Security policy to dynamically mask data at query time in *one column* based on the *value of another*

- Row level policy for column access
- No change to the stored data
- Same properties as other column level policies

```

create or replace masking policy BAZ
as (val string, optin string) returns string ->
case
  when optin = 'Y' then val
  else '***MASKED***'
end;
  
```

```

alter table start_with_a_table modify column PHONE set
masking policy BAZ using (PHONE, OPTIN);
  
```



Projection Policies

Protect values of specific columns

WHAT IS IT

Block queries that enumerate values of designated columns while allowing them in operations like filter, group, and join

WHY USE IT

Consumers can discover insights from sensitive data, while the data in designated columns is protected from exposure

HOW TO USE IT

Apply Projection Policies on sensitive columns using SQL syntax

```
CREATE PROJECTION POLICY proj_policy_ssn
AS () RETURNS PROJECTION_CONSTRAINT ->
CASE
WHEN CURRENT_ROLE() = 'ADMIN' THEN
  PROJECTION_CONSTRAINT(ALLOW => true)
ELSE
  PROJECTION_CONSTRAINT(ALLOW => false)
END;
```



```
ALTER TABLE hr.employees.directory
MODIFY COLUMN social_security_number
SET PROJECTION POLICY proj_policy_ssn;
```



```
SELECT social_security_number
FROM directory;
```



```
SELECT count(*)
FROM directory AS d JOIN customers AS
s
ON d.social_security_number =
s.social_security_number;
```



Aggregation Policies

Protect values of individual records

WHAT IS IT

Only allow aggregate queries that have more than a minimum number of rows

WHY USE IT

Consumers can discover insights from sensitive data, while individual rows are protected from exposure

HOW TO USE IT

Specify a minimum group size in a policy and then apply it to tables and views via SQL syntax

```
CREATE AGGREGATION POLICY employees_agg_policy
AS () RETURNS PROJECTION_CONSTRAINT ->
CASE
WHEN CURRENT_ROLE() = 'ADMIN' THEN
  NO_AGGREGATION_CONSTRAINT()
ELSE
  AGGREGATION_CONSTRAINT(MIN_GROUP_SIZE=>5)
END;
```



```
ALTER TABLE employees SET AGGREGATION
POLICY employees_agg_policy;
```



```
✗ SELECT * FROM employees;

✗ SELECT count(*)
    FROM employees
    WHERE name='Latanya
      Sweeney';

✓ SELECT count(*)
    FROM employees
    WHERE
      department='Sales';
```



Entity-level Privacy for Aggregation Policies

Protect values of individual entities

WHAT IS IT

Only allow aggregate queries that have more than a minimum number of entities, e.g., people, locations, organizations

WHY USE IT

Consumers can discover insights from sensitive data, while the information of individual entities across multiple rows is protected from exposure

HOW TO USE IT

Specify an optional Entity Key clause in an Aggregation Policy via SQL syntax

txn_date				txn_amount
2022-04-05	John Smith	39	79007	\$20.19
2022-04-06	John Smith	39	79007	\$13.76
2022-04-15	John Smith	39	79007	\$42.03
2022-04-21	John Smith	39	79007	\$20.19
2022-04-22	Travis Ortega	65	79010	\$61.89
⋮				

```
Query: SELECT average(age)
        FROM purchases
        WHERE name='John Smith';
```

Without entity-level privacy,
simple aggregation policy
doesn't block this sensitive query

With entity-level, the
query is blocked to
protect John Smith



Differential Privacy Policy

Protect against re-identification and privacy attacks

WHAT IS IT


The highest standard of privacy protection:
Only allow aggregate queries and add noise on sensitive data

WHY USE IT

Consumers can discover insights from granular, highly sensitive data, while protecting against re-identification and privacy attacks

HOW TO USE IT

Apply Differential Privacy Policies to tables and views via SQL syntax



Sensitive data: Full granularity, unmasked, not anonymized



Analysts can see:

Statistics & aggregates
Feature eng. for AI/ML



Analysts can't see:

See row-level data
Identify sensitive entities
Reverse-engineer proprietary info



Differential Privacy Policy: How it works

Mathematically tuned, noisy aggregates and privacy budget tracking

KEY ELEMENTS OF DIFFERENTIAL PRIVACY

NOISE

Noise is dynamically added to queries: more noise, higher protection

PRIVACY BUDGET

Blocks privacy attacks that occur across a history of queries. Consumers are cut off from running further queries before they can learn sensitive information

